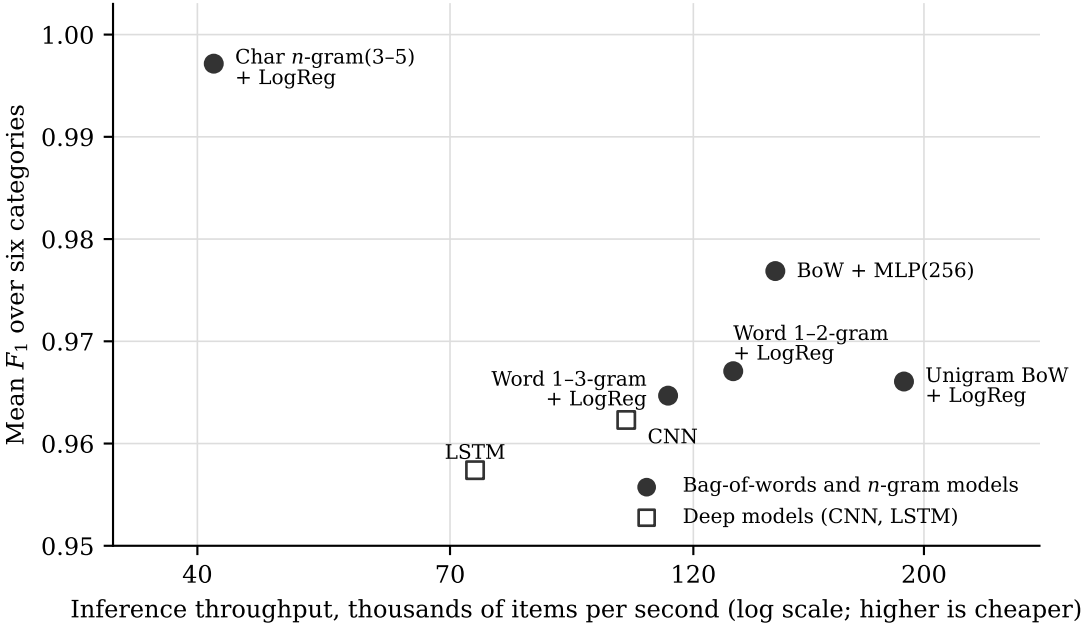
1.00
0.99
0.98
0.97
0.96
0.95
Mean F_1 over six categories
Char *n*-gram(3–5)
+ LogReg
BoW + MLP(256)
Word 1–2-gram
+ LogReg
Word 1–3-gram
+ LogReg
Unigram BoW
+ LogReg
CNN
LSTM
Bag-of-words and *n*-gram models
Deep models (CNN, LSTM)
40
70
120
200
Inference throughput, thousands of items per second (log scale; higher is cheaper)